

CS 332/532 Systems Programming

Lecture 22

Processes in OS


Professor : Mahmut Unan – UAB CS

Agenda


- Deadlock
- Signals
- Pipes

Deadlock

- The *permanent* blocking of a set of processes that either compete for system resources or communicate with each other
- A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set
- Permanent because none of the events is ever triggered
- No efficient solution in the general case



(a) Deadlock possible



(b) Deadlock

Figure 6.1 Illustration of Deadlock

INTERVIEWER: EXPLAIN DEADLOCK AND I'LL HIRE YOU




@thecrazyprogrammer

PROGRAMMER: HIRE ME AND I'LL EXPLAIN IT TO YOU

Resource Categories

Reusable

- Can be safely used by only one process at a time and is not depleted by that use
- Processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases, and semaphores




Consumable

- One that can be created (produced) and destroyed (consumed)
 - Interrupts, signals, messages, and information
 - In I/O buffers

Example 2: Memory Request


- Space is available for allocation of 200Kbytes, and the following sequence of events occur:



- Deadlock occurs if both processes progress to their second request

Consumable Resources Deadlock

- Consider a pair of processes, in which each process attempts to receive a message from the other process and then send a message to the other process:



Deadlock Approaches

- There is no single effective strategy that can deal with all types of deadlock
- Three approaches are common:
 - **Deadlock prevention**
 - Disallow one of the three necessary conditions for deadlock occurrence, or prevent circular wait condition from happening
 - **Deadlock avoidance**
 - Do not grant a resource request if this allocation might lead to deadlock
 - **Deadlock detection**
 - Grant resource requests when possible, but periodically check for the presence of deadlock and take action to recover

Conditions for Deadlock

Mutual Exclusion

- Only one process may use a resource at a time
- No process may access a resource until that has been allocated to another process

Hold-and-Wait

- A process may hold allocated resources while awaiting assignment of other resources

No Pre-emption

- No resource can be forcibly removed from a process holding it

Circular Wait


- A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain

Deadlock Prevention Strategy

- Design a system in such a way that the possibility of deadlock is excluded
- Two main methods:
 - Indirect
 - Prevent the occurrence of one of the three necessary conditions
 - Direct
 - Prevent the occurrence of a circular wait

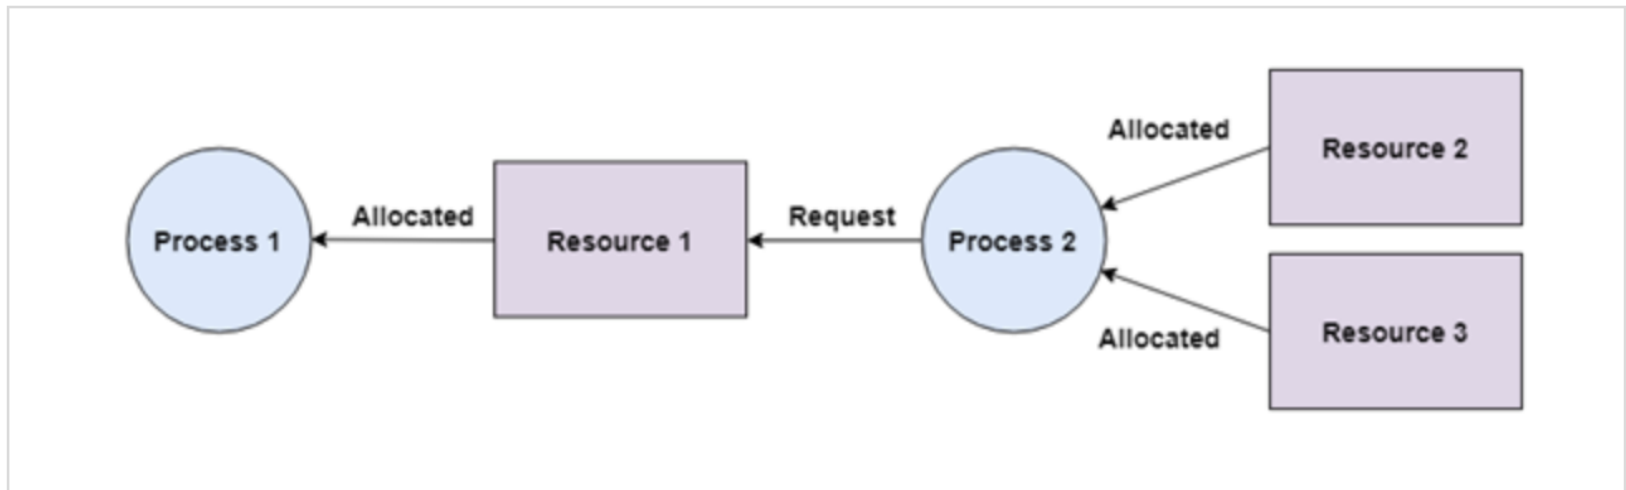
Deadlock Condition Prevention

- Mutual exclusion
 - If access to a resource requires mutual exclusion, then mutual exclusion must be supported by the OS
 - Some resources, such as files, may allow multiple accesses for reads but only exclusive access for writes
 - Even in this case, deadlock can occur if more than one process requires write permission



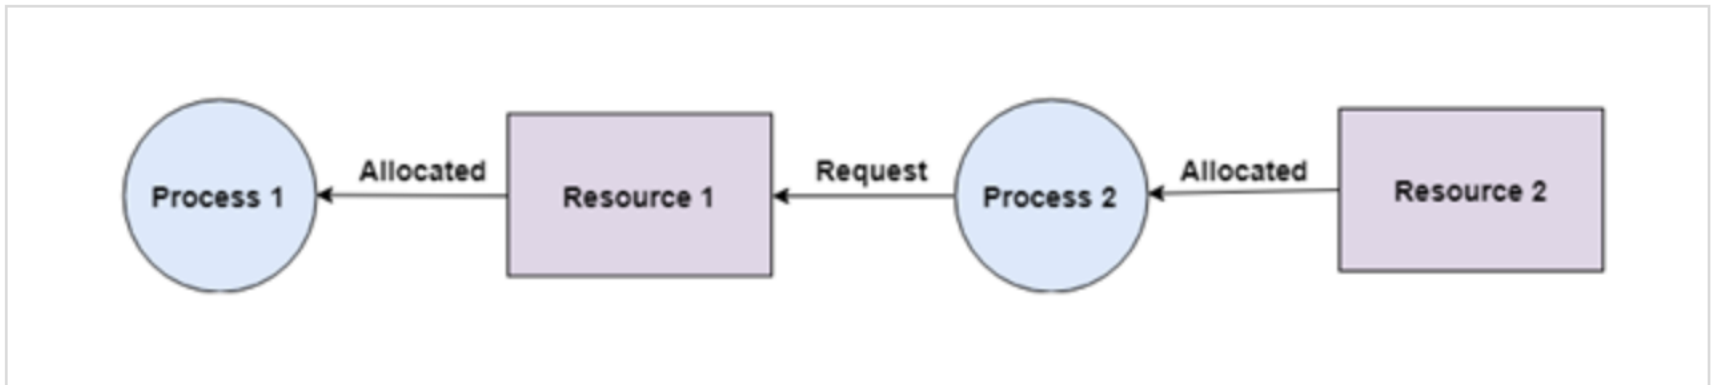
Deadlock Condition Prevention

- Hold and wait
 - Can be prevented by requiring that a process request all of its required resources at one time and blocking the process until all requests can be granted simultaneously



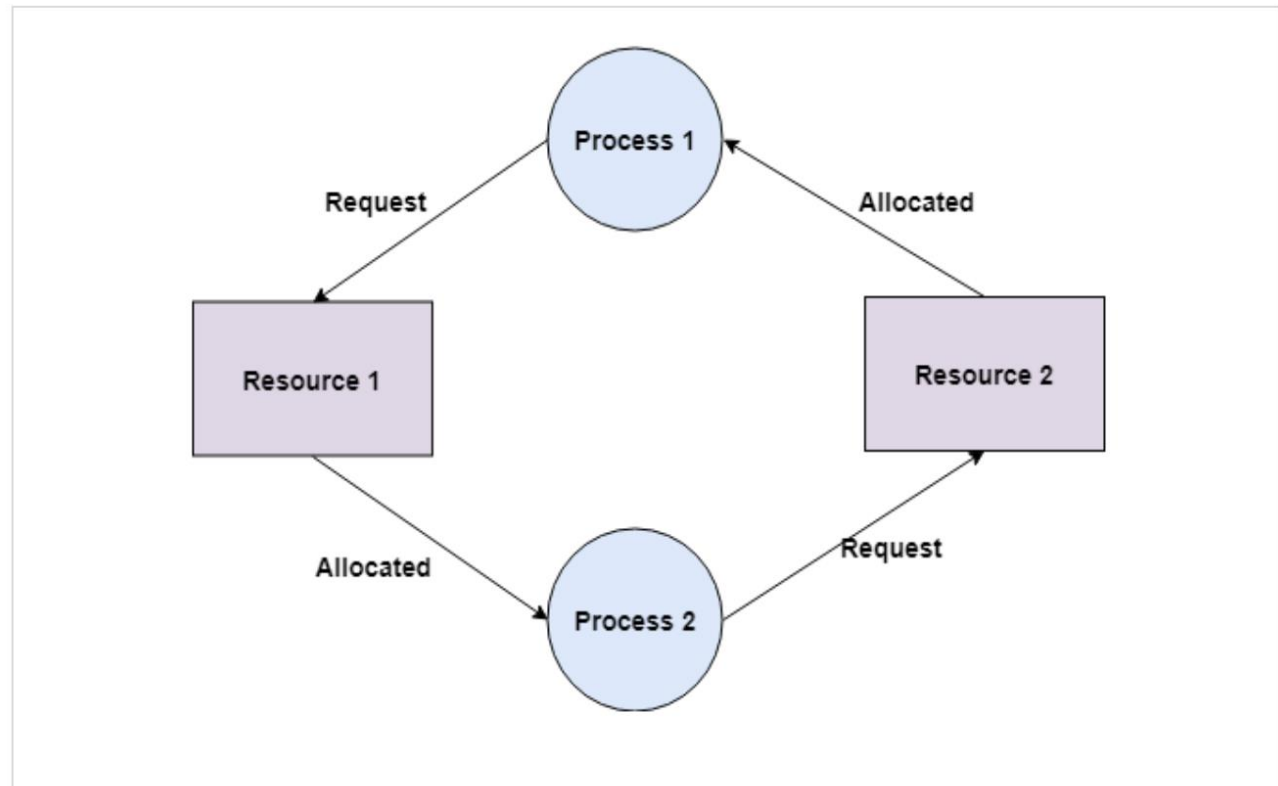
Deadlock Condition Prevention

- No Preemption
 - If a process holding certain resources is denied a further request, that process must release its original resources and request them again
 - OS may preempt the second process and require it to release its resources



Deadlock Condition Prevention

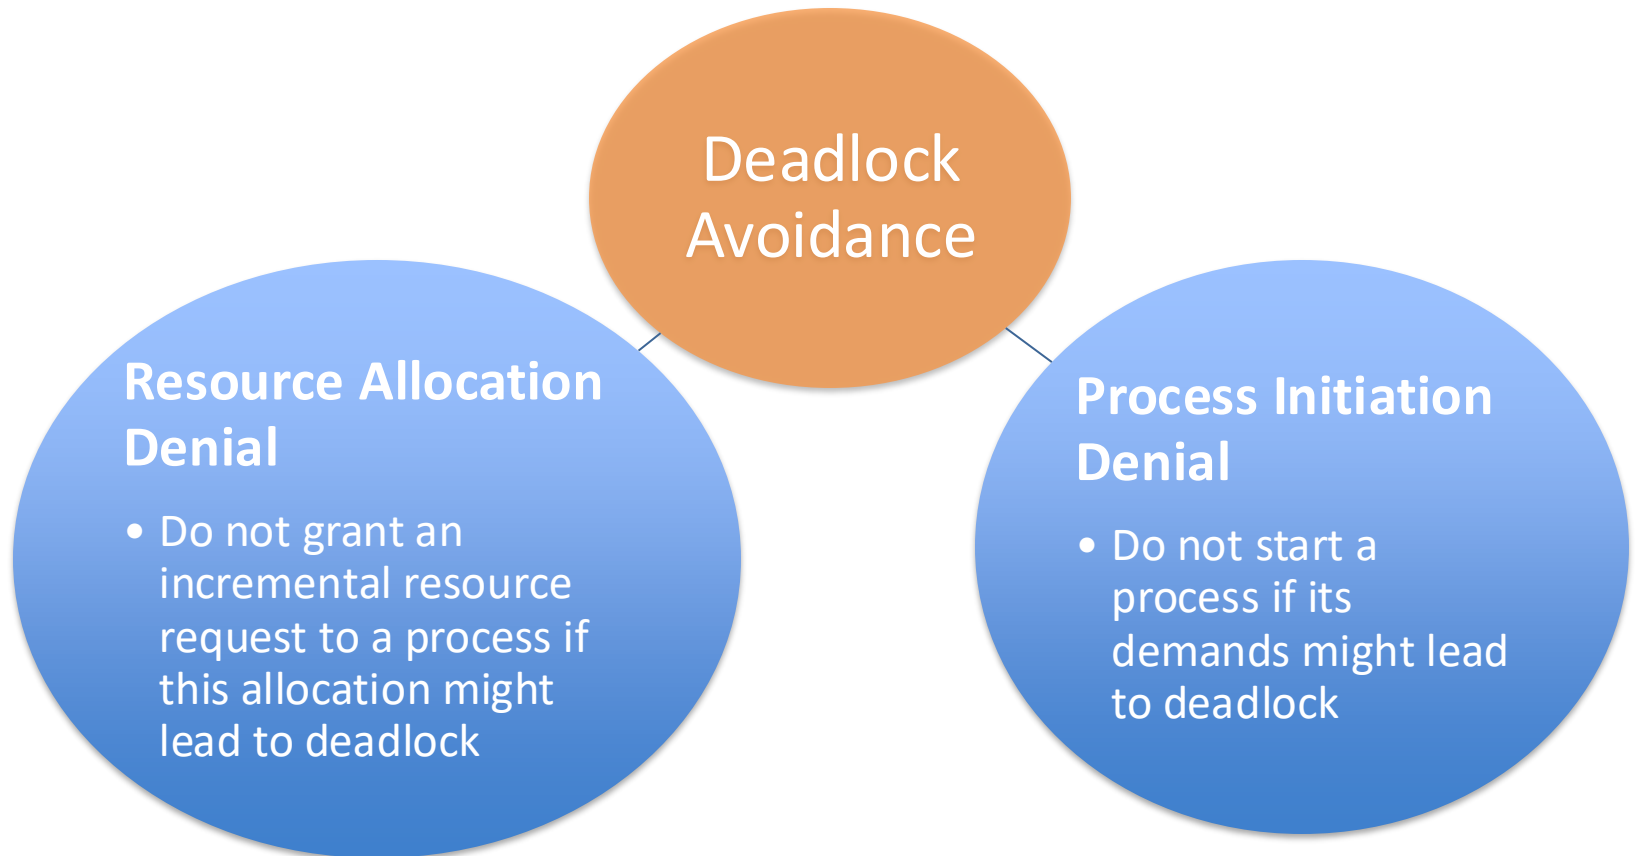
- Circular Wait
 - The circular wait condition can be prevented by defining a linear ordering of resource types



Deadlock Avoidance

- Allows the three necessary conditions but makes judicious choices to assure that the deadlock point is never reached
- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Allows the three necessary conditions but makes judicious choices to assure that the deadlock point is never reached
- Requires knowledge of future process requests

Two Approaches to Deadlock Avoidance




Resource Allocation Denial

- Referred to as the *banker's algorithm*
- **State** of the system reflects the current allocation of resources to processes
- **Safe state** is one in which there is at least one sequence of resource allocations to processes that does not result in a deadlock
- **Unsafe state** is a state that is not safe

Deadlock Avoidance Advantages

- It is not necessary to preempt and rollback processes, as in deadlock detection
- It is less restrictive than deadlock prevention

Deadlock Avoidance Restrictions

- 
- Maximum resource requirement for each process must be stated in advance
 - Processes under consideration must be independent and with no synchronization requirements
 - There must be a fixed number of resources to allocate
 - No process may exit while holding resources


Deadlock Detection Algorithm

- A check for deadlock can be made as frequently as each resource request or, less frequently, depending on how likely it is for a deadlock to occur



Advantages:

- It leads to early detection
- The algorithm is relatively simple



Disadvantage

- Frequent checks consume considerable processor time

UNIX Concurrency Mechanisms

- UNIX provides a variety of mechanisms for interprocessor communication and synchronization including:

Pipes

Messages

Shared
memory

Semaphores

Signals

Pipes

- Circular buffers allowing two processes to communicate on the producer-consumer model
 - First-in-first-out queue, written by one process and read by another

Two types:

- Named
- Unnamed

Messages

- A block of bytes with an accompanying type
- UNIX provides ***msgsnd*** and ***msgrcv*** system calls for processes to engage in message passing
- Associated with each process is a message queue, which functions like a mailbox

Shared Memory

- Fastest form of interprocess communication
- Common block of virtual memory shared by multiple processes
- Permission is read-only or read-write for a process
- Mutual exclusion constraints are not part of the shared-memory facility but must be provided by the processes using the shared memory

Semaphores

- Generalization of the `semWait` and `semSignal` primitives
 - No other process may access the semaphore until all operations have completed

Consists of:

- Current value of the semaphore
- Process ID of the last process to operate on the semaphore
- Number of processes waiting for the semaphore value to be greater than its current value
- Number of processes waiting for the semaphore value to be zero

Signals

- A software mechanism that informs a process of the occurrence of asynchronous events
 - Similar to a hardware interrupt, but does not employ priorities
- A signal is delivered by updating a field in the process table for the process to which the signal is being sent
- A process may respond to a signal by:
 - Performing some default action
 - Executing a signal-handler function
 - Ignoring the signal

Value	Name	Description
01	SIGHUP	Hang up; sent to process when kernel assumes that the user of that process is doing no useful work
02	SIGINT	Interrupt
03	SIGQUIT	Quit; sent by user to induce halting of process and production of core dump
04	SIGILL	Illegal instruction
05	SIGTRAP	Trace trap; triggers the execution of code for process tracing
06	SIGIOT	IOT instruction
07	SIGEMT	EMT instruction
08	SIGFPE	Floating-point exception
09	SIGKILL	Kill; terminate process
10	SIGBUS	Bus error
11	SIGSEGV	Segmentation violation; process attempts to access location outside its virtual address space
12	SIGSYS	Bad argument to system call
13	SIGPIPE	Write on a pipe that has no readers attached to it
14	SIGALRM	Alarm clock; issued when a process wishes to receive a signal after a period of time
15	SIGTERM	Software termination
16	SIGUSR1	User-defined signal 1
17	SIGUSR2	User-defined signal 2
18	SIGCHLD	Death of a child
19	SIGPWR	Power failure

Table 6.2

UNIX Signals

(Table can be found on page 288 in textbook)

Real-time (RT) Signals

- Linux includes all of the concurrency mechanisms found in other UNIX systems
- Linux also supports real-time (RT) signals
- RT signals differ from standard UNIX signals in three primary ways:
 - Signal delivery in priority order is supported
 - Multiple signals can be queued
 - With standard signals, no value or message can be sent to the target process – it is only a notification
 - With RT signals it is possible to send a value along with the signal

Atomic Operations

- Atomic operations execute without interruption and without interference
- Simplest of the approaches to kernel synchronization
- Two types:

Integer Operations

Operate on an integer variable

Typically used to implement counters

Bitmap Operations

Operate on one of a sequence of bits at an arbitrary memory location indicated by a pointer variable

Atomic Integer Operations	
ATOMIC_INIT (int i)	At declaration: initialize an atomic_t to i
int atomic_read(atomic_t *v)	Read integer value of v
void atomic_set(atomic_t *v, int i)	Set the value of v to integer i
void atomic_add(int i, atomic_t *v)	Add i to v
void atomic_sub(int i, atomic_t *v)	Subtract i from v
void atomic_inc(atomic_t *v)	Add 1 to v
void atomic_dec(atomic_t *v)	Subtract 1 from v
int atomic_sub_and_test(int i, atomic_t *v)	Subtract i from v; return 1 if the result is zero; return 0 otherwise
int atomic_add_negative(int i, atomic_t *v)	Add i to v; return 1 if the result is negative; return 0 otherwise (used for implementing semaphores)
int atomic_dec_and_test(atomic_t *v)	Subtract 1 from v; return 1 if the result is zero; return 0 otherwise
int atomic_inc_and_test(atomic_t *v)	Add 1 to v; return 1 if the result is zero; return 0 otherwise
Atomic Bitmap Operations	
void set_bit(int nr, void *addr)	Set bit nr in the bitmap pointed to by addr
void clear_bit(int nr, void *addr)	Clear bit nr in the bitmap pointed to by addr
void change_bit(int nr, void *addr)	Invert bit nr in the bitmap pointed to by addr
int test_and_set_bit(int nr, void *addr)	Set bit nr in the bitmap pointed to by addr; return the old bit value
int test_and_clear_bit(int nr, void *addr)	Clear bit nr in the bitmap pointed to by addr; return the old bit value
int test_and_change_bit(int nr, void *addr)	Invert bit nr in the bitmap pointed to by addr; return the old bit value
int test_bit(int nr, void *addr)	Return the value of bit nr in the bitmap pointed to by addr

Table 6.2

Linux Atomic Operations

(Table can be found on page 289 in textbook)

Spinlocks

- Most common technique for protecting a critical section in Linux
- Can only be acquired by one thread at a time
 - Any other thread will keep trying (spinning) until it can acquire the lock
- Built on an integer location in memory that is checked by each thread before it enters its critical section
- Effective in situations where the wait time for acquiring a lock is expected to be very short
- Disadvantage:
 - Locked-out threads continue to execute in a busy-waiting mode

<code>void spin_lock(spinlock_t *lock)</code>	Acquires the specified lock, spinning if needed until it is available
<code>void spin_lock_irq(spinlock_t *lock)</code>	Like spin lock, but also disables interrupts on the local processor
<code>void spin_lock_irqsave(spinlock_t *lock, unsigned long flags)</code>	Like spin lock irq, but also saves the current interrupt state in flags
<code>void spin_lock_bh(spinlock_t *lock)</code>	Like spin lock, but also disables the execution of all bottom halves
<code>void spin_unlock(spinlock_t *lock)</code>	Releases given lock
<code>void spin_unlock_irq(spinlock_t *lock)</code>	Releases given lock and enables local interrupts
<code>void spin_unlock_irqrestore(spinlock_t *lock, unsigned long flags)</code>	Releases given lock and restores local interrupts to given previous state
<code>void spin_unlock_bh(spinlock_t *lock)</code>	Releases given lock and enables bottom halves
<code>void spin_lock_init(spinlock_t *lock)</code>	Initializes given spinlock
<code>int spin_trylock(spinlock_t *lock)</code>	Tries to acquire specified lock; returns nonzero if lock is currently held and zero otherwise
<code>int spin_is_locked(spinlock_t *lock)</code>	Returns nonzero if lock is currently held and zero otherwise

Table 6.4 Linux Spinlocks

(Table can be found on page 291 in textbook)

Semaphores

- User level:
 - Linux provides a semaphore interface corresponding to that in UNIX SVR4
- Internally:
 - Implemented as functions within the kernel and are more efficient than user-visible semaphores
- Three types of kernel semaphores:
 - Binary semaphores
 - Counting semaphores
 - Reader-writer semaphores

Table 6.5

Linux

Semaphores

Traditional Semaphores	
<code>void sema_init(struct semaphore *sem, int count)</code>	Initializes the dynamically created semaphore to the given count
<code>void init MUTEX(struct semaphore *sem)</code>	Initializes the dynamically created semaphore with a count of 1 (initially unlocked)
<code>void init MUTEX LOCKED(struct semaphore *sem)</code>	Initializes the dynamically created semaphore with a count of 0 (initially locked)
<code>void down(struct semaphore *sem)</code>	Attempts to acquire the given semaphore, entering uninterruptible sleep if semaphore is unavailable
<code>int down interruptible(struct semaphore *sem)</code>	Attempts to acquire the given semaphore, entering interruptible sleep if semaphore is unavailable; returns <code>-EINTR</code> value if a signal other than the result of an up operation is received
<code>int down trylock(struct semaphore *sem)</code>	Attempts to acquire the given semaphore, and returns a nonzero value if semaphore is unavailable
<code>void up(struct semaphore *sem)</code>	Releases the given semaphore
Reader-Writer Semaphores	
<code>void init rwsem(struct rw_semaphore, *rwsem)</code>	Initializes the dynamically created semaphore with a count of 1
<code>void down read(struct rw semaphore, *rwsem)</code>	Down operation for readers
<code>void up read(struct rw semaphore, *rwsem)</code>	Up operation for readers
<code>void down write(struct rw_semaphore, *rwsem)</code>	Down operation for writers
<code>void up write(struct rw semaphore, *rwsem)</code>	Up operation for writers

(Table can be found on page 293 in textbook)

Table 6.6

Linux Memory Barrier Operations

<code>rmb()</code>	Prevents loads from being reordered across the barrier
<code>wmb()</code>	Prevents stores from being reordered across the barrier
<code>mb()</code>	Prevents loads and stores from being reordered across the barrier
<code>Barrier()</code>	Prevents the compiler from reordering loads or stores across the barrier
<code>smp_rmb()</code>	On SMP, provides a <code>rmb()</code> and on UP provides a <code>barrier()</code>
<code>smp_wmb()</code>	On SMP, provides a <code>wmb()</code> and on UP provides a <code>barrier()</code>
<code>smp_mb()</code>	On SMP, provides a <code>mb()</code> and on UP provides a <code>barrier()</code>

SMP = symmetric multiprocessor

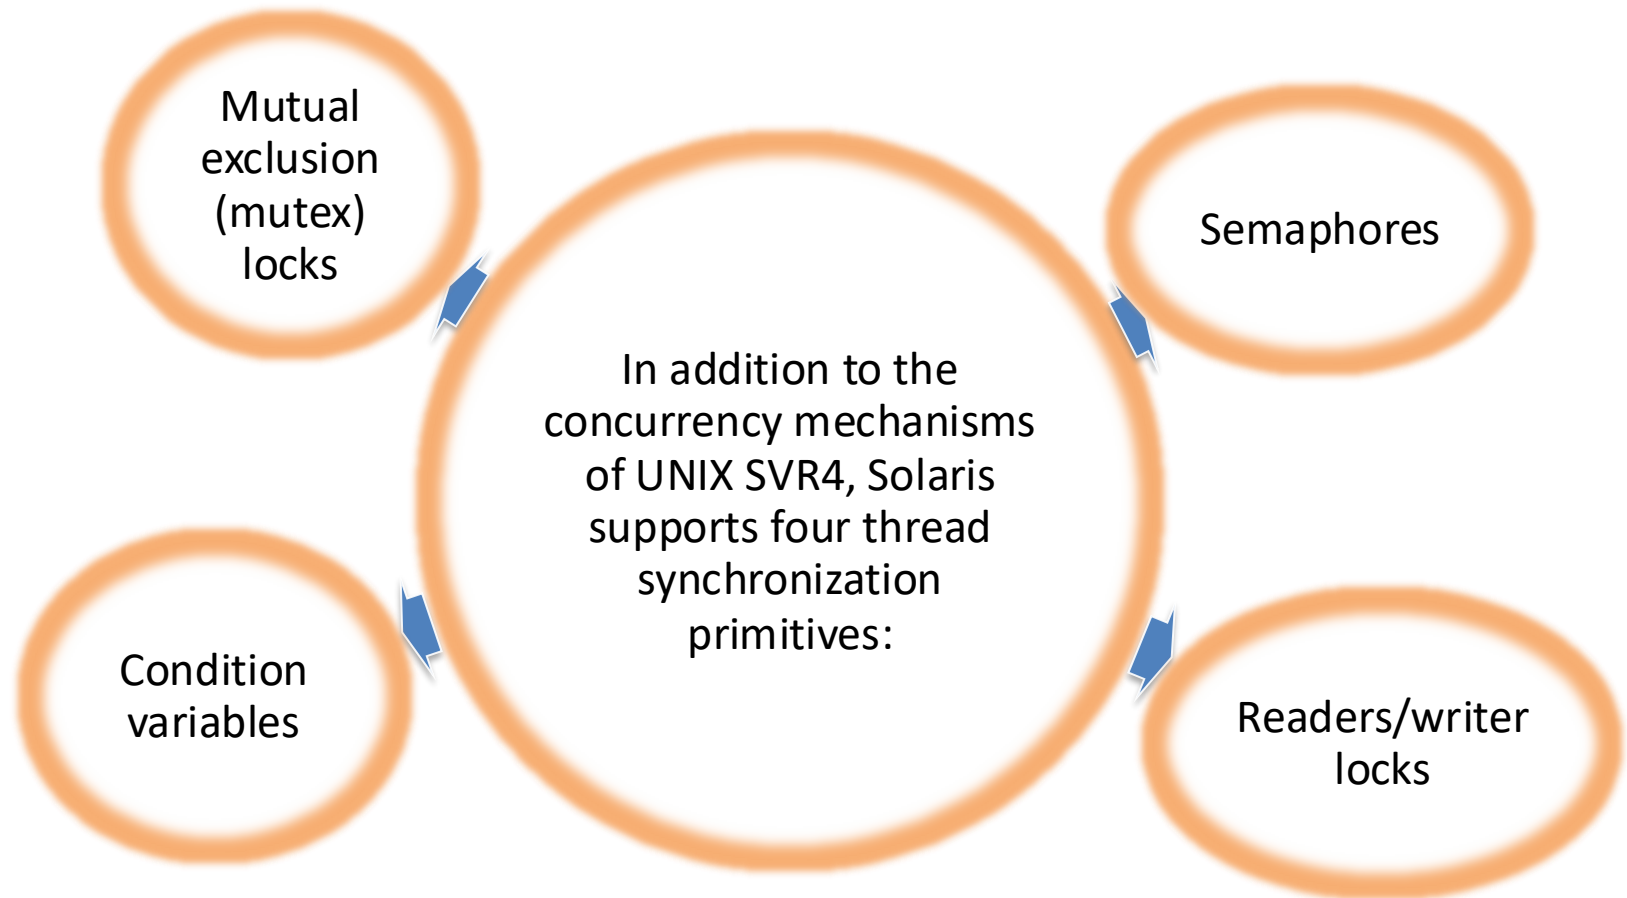
UP = uniprocessor

(Table can be found on page 294 in textbook)

Read-Copy-Update (RCU)

- The RCU mechanism is an advanced lightweight synchronization mechanism which was integrated into the Linux kernel in 2002
- The RCU is used widely in the Linux kernel
- RCU is also used by other operating systems
- There is a userspace RCU library called liburcu
- The shared resources that the RCU mechanism protects must be accessed via a pointer
- The RCU mechanism provides access for multiple readers and writers to a shared resource

Synchronization Primitives



Mutual Exclusion (MUTEX) Lock

- Used to ensure only one thread at a time can access the resource protected by the mutex
- The thread that locks the mutex must be the one that unlocks it
- A thread attempts to acquire a mutex lock by executing the `mutex_enter` primitive
- Default blocking policy is a spinlock
- An interrupt-based blocking mechanism is optional

Semaphores

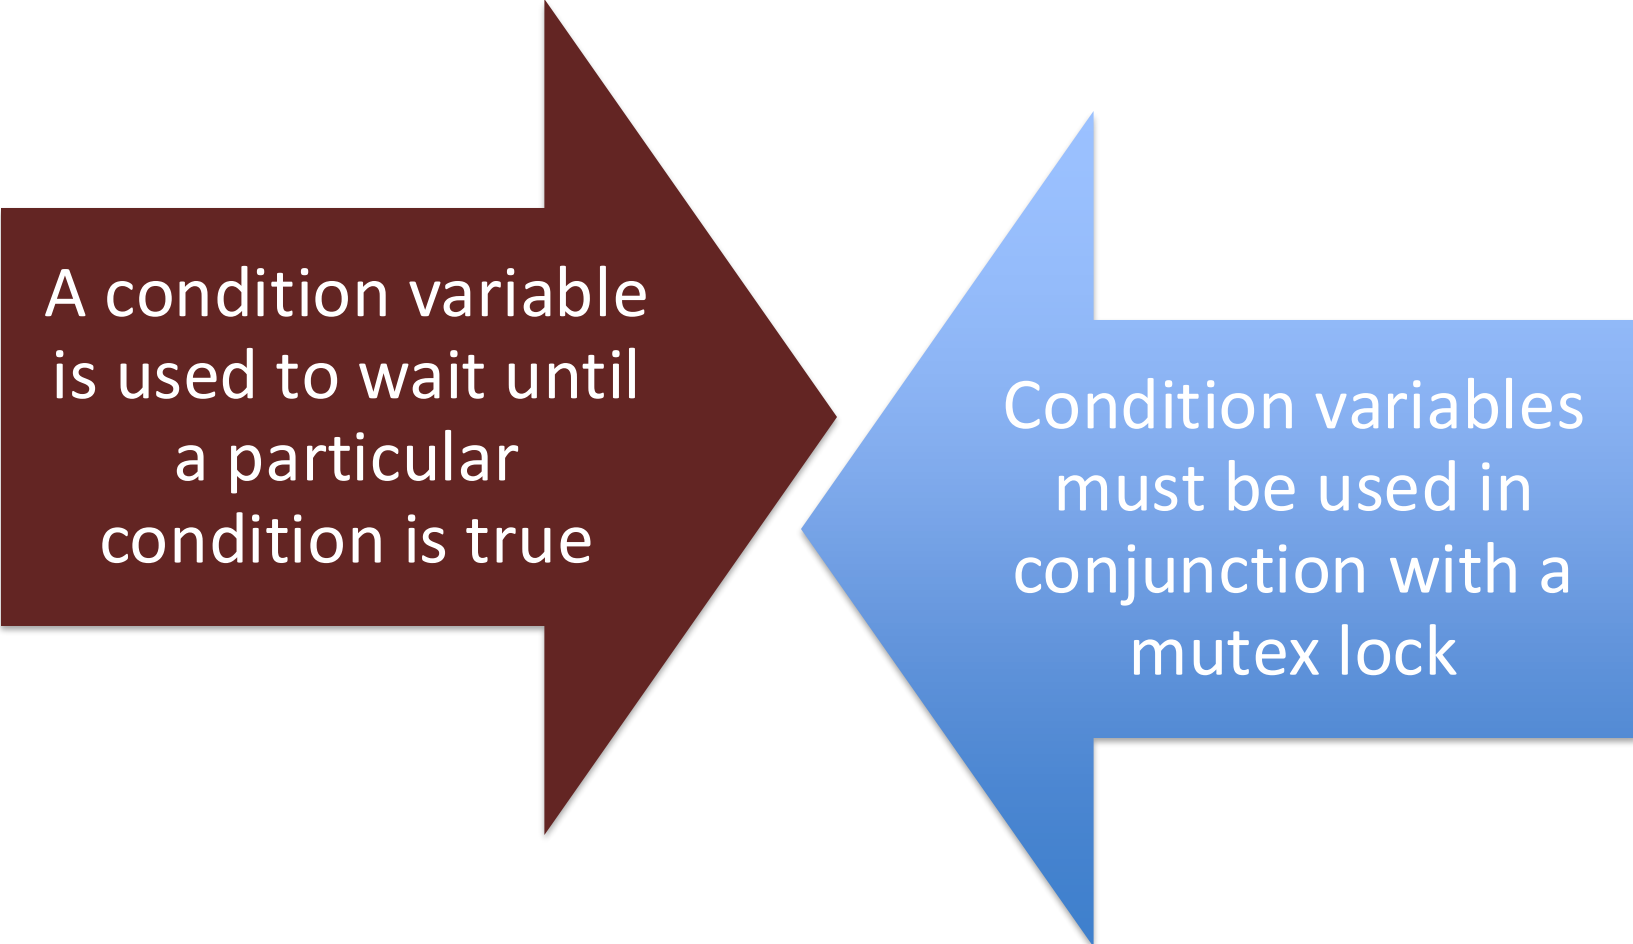
Solaris provides classic counting semaphores with the following primitives:

- `sema_p()` Decrements the semaphore, potentially blocking the thread
- `sema_v()` Increments the semaphore, potentially unblocking a waiting thread
- `sema_tryv()` Decrements the semaphore if blocking is not required

Readers/Writer Locks

- Allows multiple threads to have simultaneous read-only access to an object protected by the lock
- Allows a single thread to access the object for writing at one time, while excluding all readers
 - When lock is acquired for writing it takes on the status of `write lock`
 - If one or more readers have acquired the lock its status is `read lock`

Condition Variables



A condition variable
is used to wait until
a particular
condition is true

Condition variables
must be used in
conjunction with a
mutex lock