

# CS 332/532 Systems Programming

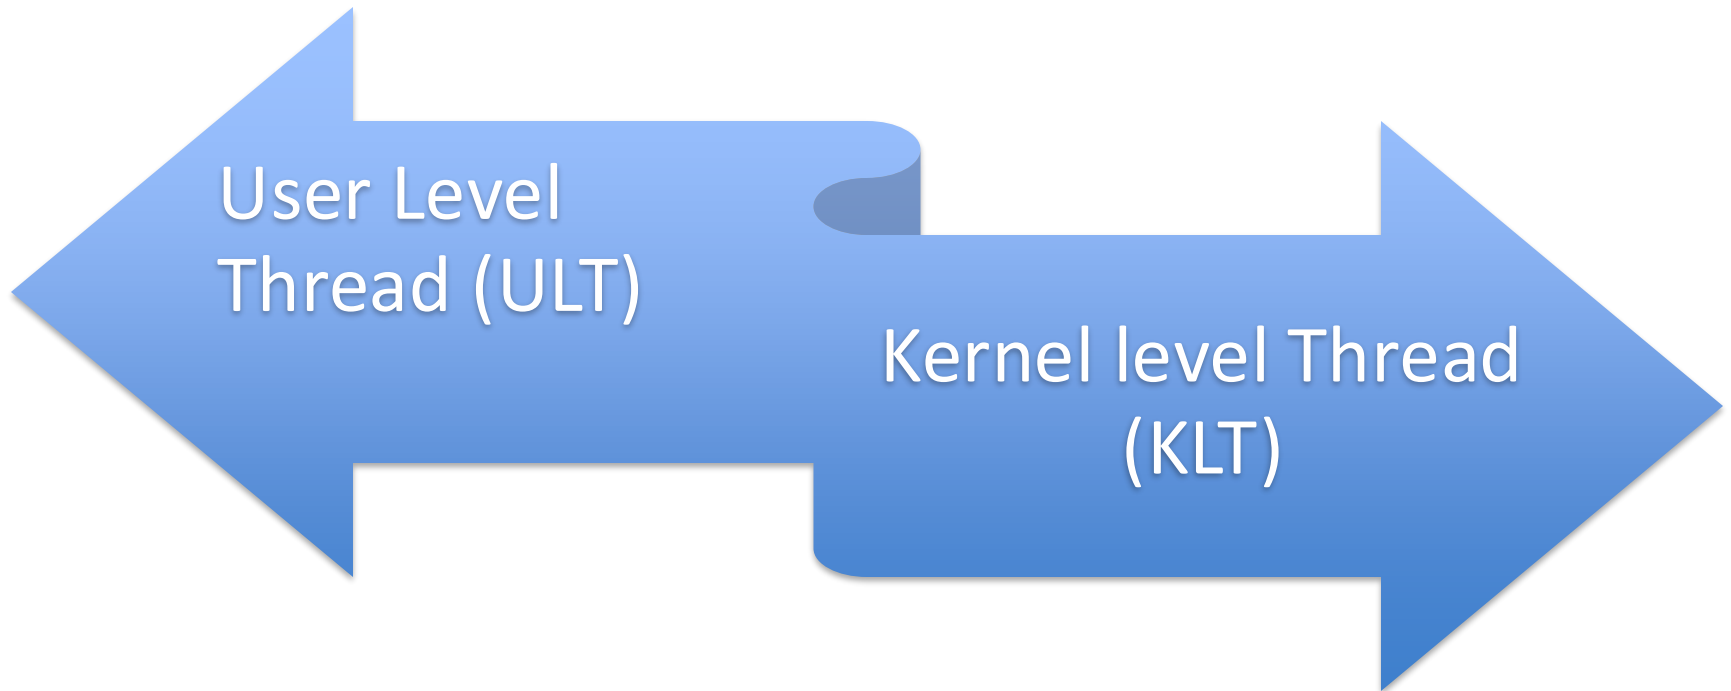
## Lecture 29 Threads

Professor : Mahmut Unan – UAB CS

# Agenda

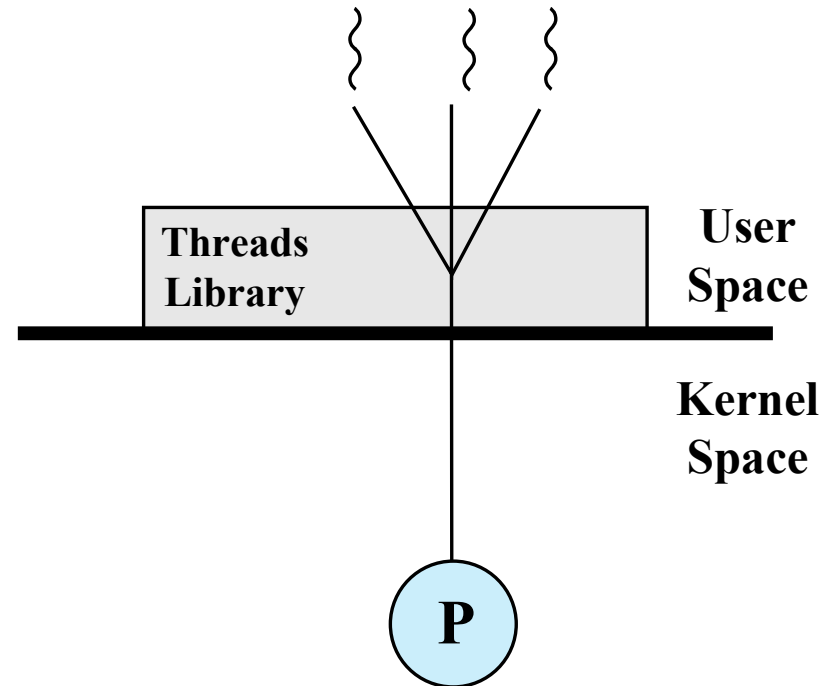
- Threads

# Types of Threads

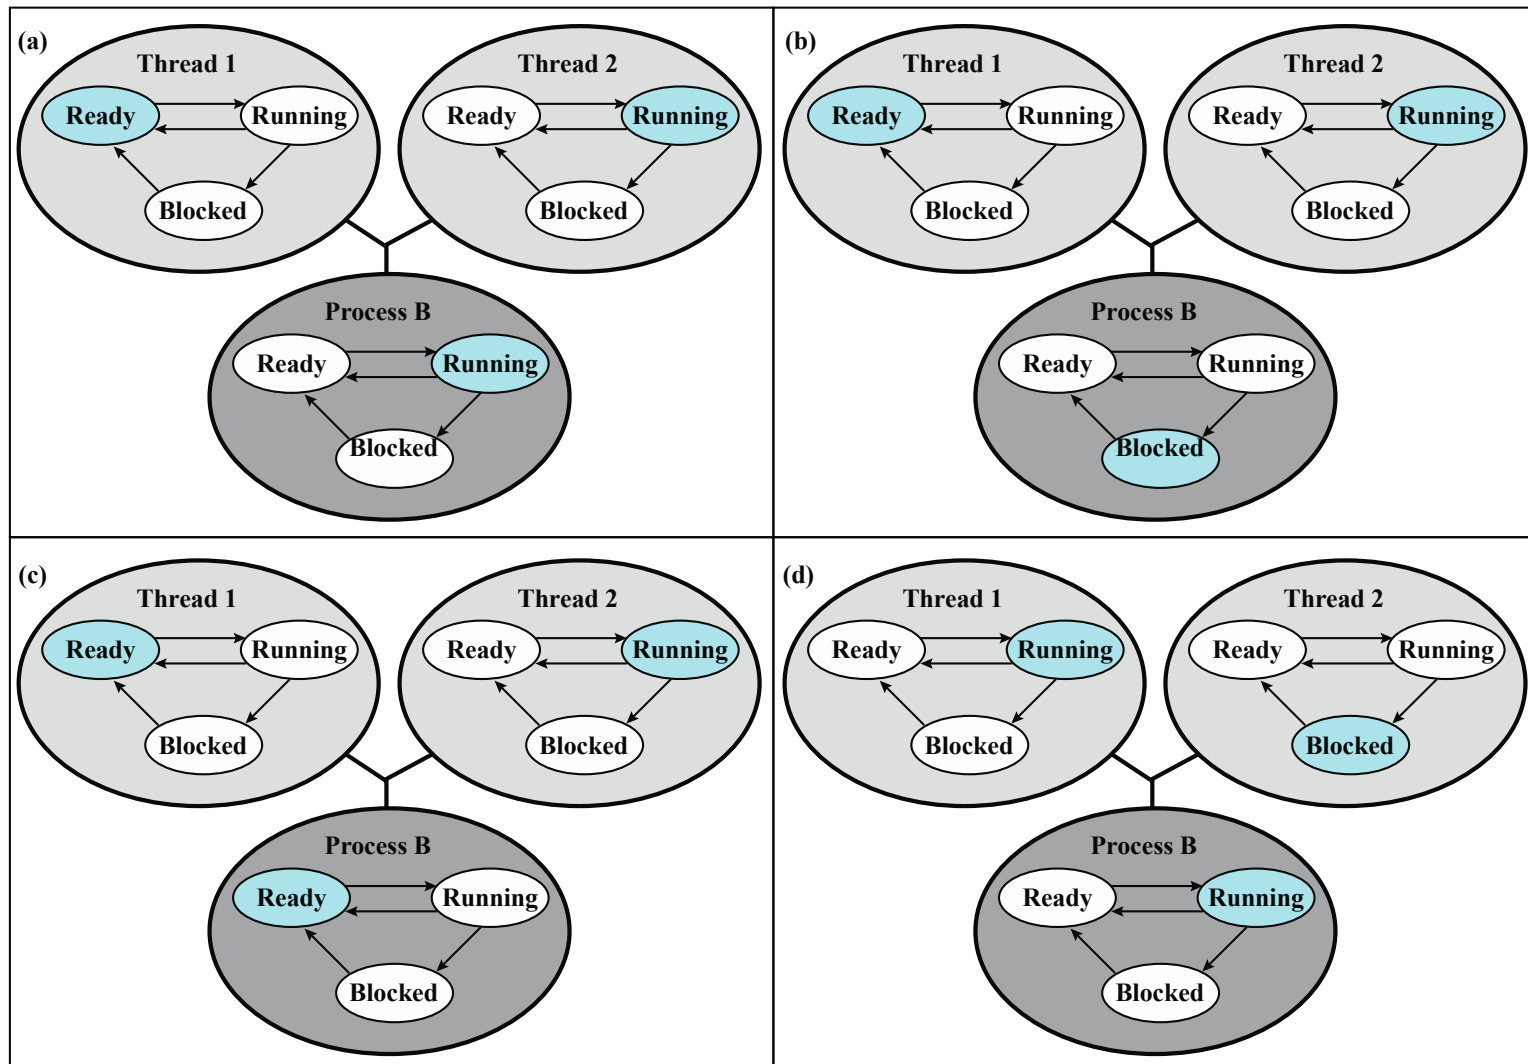


# User-Level Threads (ULTs)

- All thread management is done by the application
- The kernel is not aware of the existence of threads



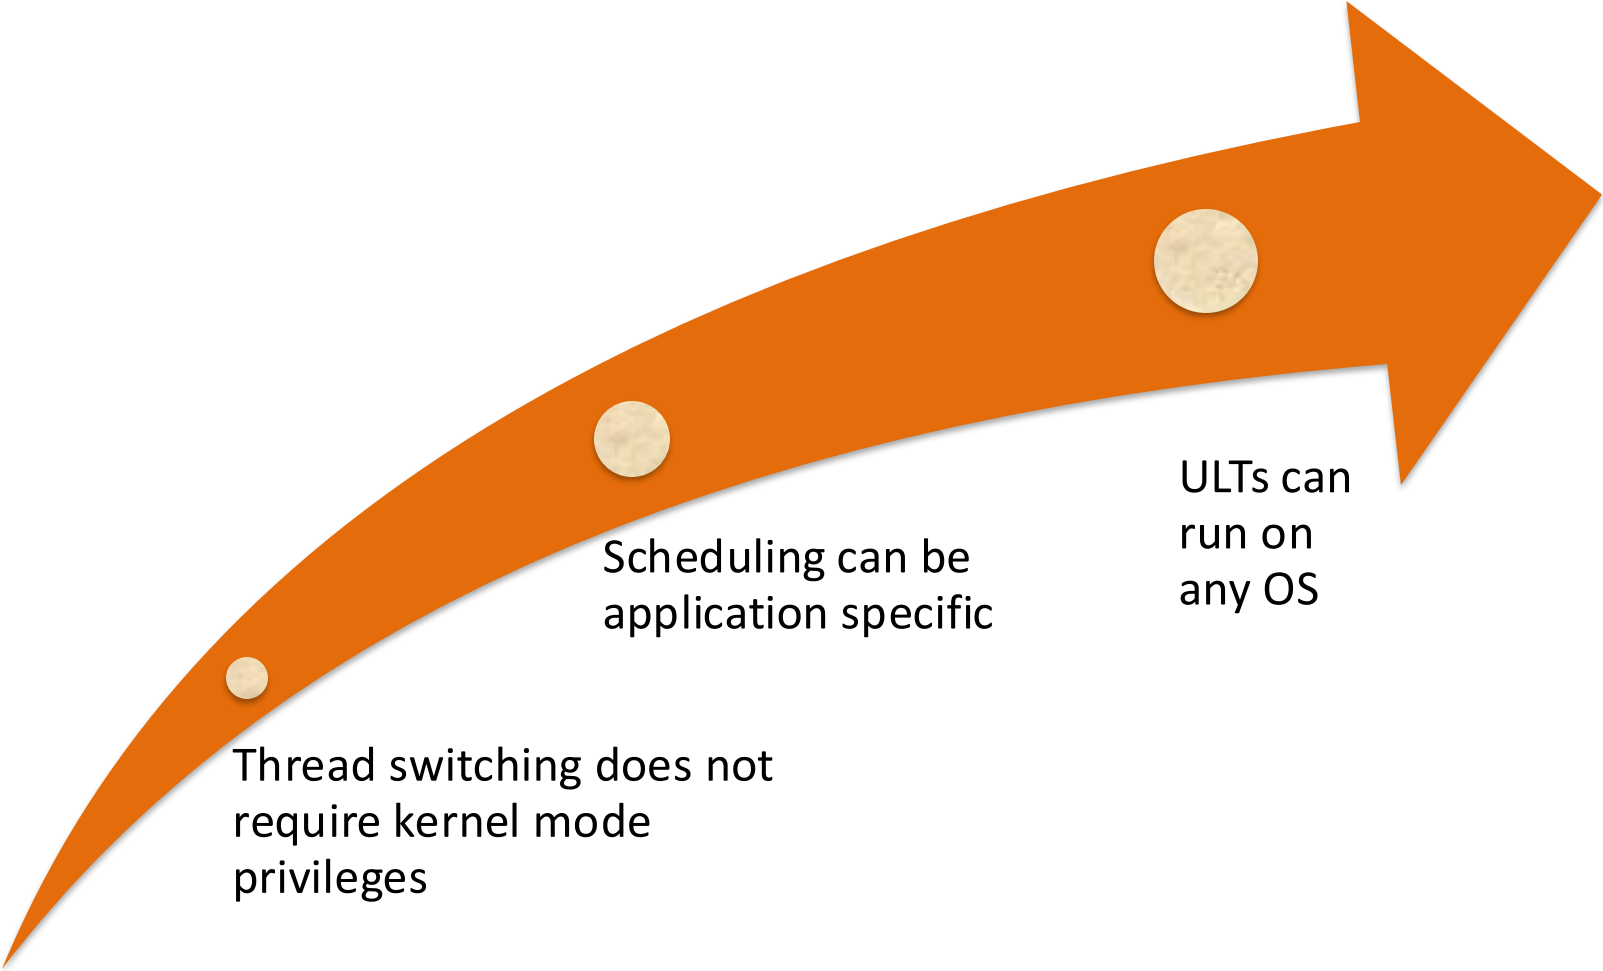
**(a) Pure user-level**



Colored state  
is current state

**Figure 4.6 Examples of the Relationships Between User-Level Thread States and Process States**

# Advantages of ULTs



Thread switching does not  
require kernel mode  
privileges

Scheduling can be  
application specific

ULTs can  
run on  
any OS

# Disadvantages of ULTs

- In a typical OS many system calls are blocking
  - As a result, when a ULT executes a system call, not only is that thread blocked, but all of the threads within the process are blocked as well
- In a pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing
  - A kernel assigns one process to only one processor at a time, therefore, only a single thread within a process can execute at a time

# Overcoming ULT Disadvantages

## Jacketing

- Purpose is to convert a blocking system call into a non-blocking system call

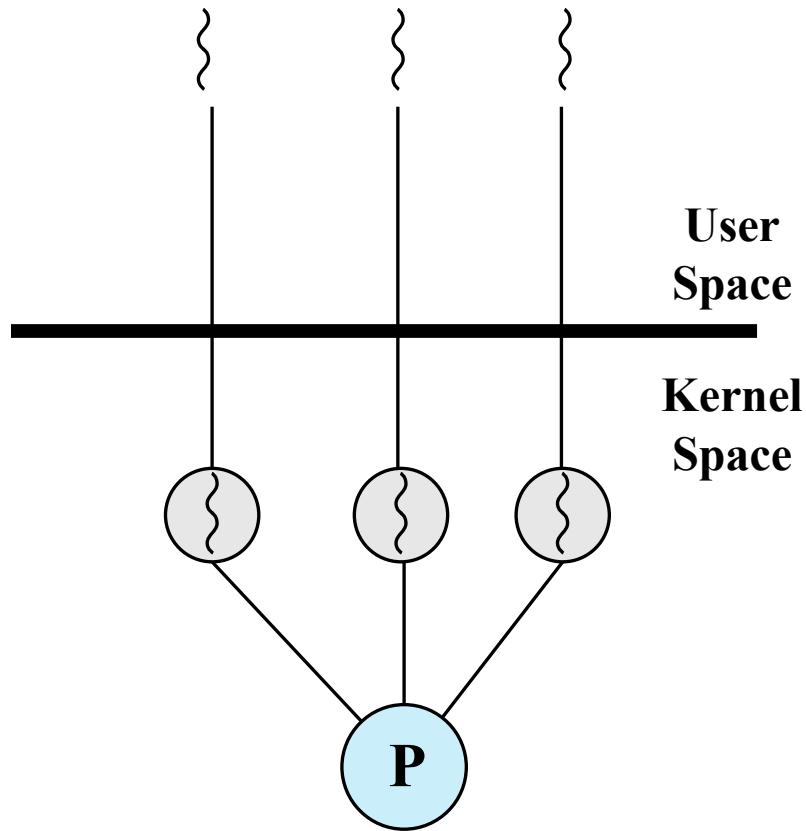


Writing an application as multiple processes rather than multiple threads

- However, this approach eliminates the main advantage of threads



# Kernel-Level Threads (KLTs)



**(b) Pure kernel-level**

- Thread management is done by the kernel
  - There is no thread management code in the application level, simply an application programming interface (API) to the kernel thread facility
  - Windows is an example of this approach

# Advantages of KLTs

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines themselves can be multithreaded

# Disadvantage of KLTs

- The transfer of control from one thread to another within the same process requires a mode switch to the kernel

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

**Table 4.1**  
**Thread and Process Operation Latencies ( $\mu$ s)**

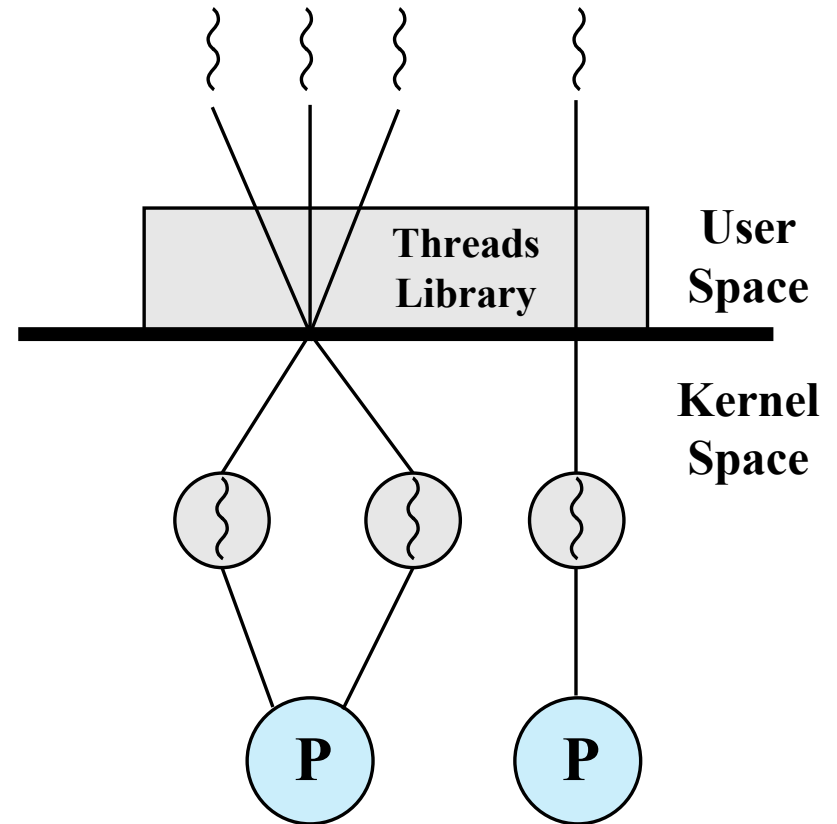
S.N.	User-Level Threads	Kernel-Level Thread
1	User-level threads are faster to create and manage.	Kernel-level threads are slower to create and manage.
2	Implementation is by a thread library at the user level.	Operating system supports creation of Kernel threads.
3	User-level thread is generic and can run on any operating system.	Kernel-level thread is specific to the operating system.
4	Multi-threaded applications cannot take advantage of multiprocessing.	Kernel routines themselves can be multithreaded.

# Multithreading Models

- Many to many relationship.
- Many to one relationship.
- One to one relationship.

# Combined Approaches

- Thread creation is done completely in the user space, as is the bulk of the scheduling and synchronization of threads within an application
- Solaris is a good example



(c) Combined

<b>Threads:Processes</b>	<b>Description</b>	<b>Example Systems</b>
<b>1:1</b>	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
<b>M:1</b>	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
<b>1:M</b>	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
<b>M:N</b>	Combines attributes of M:1 and 1:M cases.	TRIX

**Table 4.2**  
**Relationship between Threads and Processes**

# Create threads using POSIX threads library

- In the previous lectures/labs we focused on how to create processes, in this lab we will focus on creating threads and mechanisms for establishing synchronization among threads.
- First, let us understand the difference between a process and a thread.
  - A process could be considered to have two characteristics:
    - (a) resource ownership
    - (b) scheduling or execution.
- The unit of scheduling and dispatching is usually referred to as a **thread** or **lightweight process** and the ability of to support multiple, concurrent paths of execution within a single process is often referred to as *multithreading*.



- Threads offer several benefits compared to a process:
  - Threads takes less time to create a new thread than a process
  - Threads take less time to terminate a thread than a process
  - Switching between two threads (context switching) takes less time than switching between processes
  - All of the threads in a process share the state and resources of that process (since threads reside in the same address space and have access to the same data)
  - Threads enhance efficiency in communication between programs (since threads share memory and files within the same process and can communicate without invoking the kernel)

- As a result of these advantages, if we have to implement a set of functions that are closely related, implementing this functionality using multiple threads is far more efficient than using multiple processes.
- We will use the **POSIX threads library**, usually referred to as Pthreads library, that provides C APIs to create and manage threads. We have to include the file *pthread.h* and link with *-lpthread* to compile and link.
- We can create new threads using the *pthread\_create()* function which has the following function definition:

```
#include <pthread.h>
int pthread_create(pthread_t *thread, const
pthread_attr_t *attr,
void *(*start_routine) (void *), void *arg);
```

- The new thread that will be created by the *pthread\_create* function will invoke the function *start\_routine*.
- Note that the function *start\_routine* takes one argument of type *void \** and has the return type as *void \**.
- In other words, the function *start\_routine* has the following function definition:

```
void *start_routine(void *arg)
```

- When the *pthread\_create* call returns successfully, it returns the thread ID associated with the new thread created in the variable *thread*.
- This can be used by the main thread in subsequent pthread function calls such as *pthread\_join*.
- The second argument, *attr*, provides a reference to the *pthread\_attr\_t* structure that describes the various attributes of the new thread to be created.
- It can be initialized using *pthread\_attr\_init* call or set to NULL if default attributes must be used.
- You can find out more about the different thread attributes that can be specified by looking at the man page for *pthread\_attr\_init*.

- The new thread created will terminate when the function *start\_routine* returns or when a call to *pthread\_exit* is made inside the *start\_routine*.
- We can use the *pthread\_join* function to wait for a thread to complete using the thread ID that was returned when *pthread\_create* call was invoked.
- If a thread has already completed,
  - *pthread\_join* will return immediately, otherwise, it will wait for the corresponding thread to complete.

# exercise 1

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <pthread.h>
5
6
7  void *someFuncToCreateThread(void *someValue)
8  {
9      sleep(2);
10     printf("I am inside the thread \n");
11     return NULL;
12 }
13
14 int main()
15 {
16     pthread_t thread_id;
17     printf("I am inside the main function\n");
18     pthread_create(&thread_id, NULL, someFuncToCreateThread, NULL);
19     pthread_join(thread_id, NULL);
20     printf("Back to the main function\n");
21     exit(0);
22 }
23
```

# compile & run

To compile a multithreaded program, we will be using gcc and we need to link it with the pthreads library.

```
(base) mahmutunan@MacBook-Pro lecture31 % gcc exercise1.c -o exercise1 -lpthread
(base) mahmutunan@MacBook-Pro lecture31 % ./exercise1
I am inside the main function
I am inside the thread
Back to the main function
(base) mahmutunan@MacBook-Pro lecture31 % _
```

# exercise 2

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

void *function1(void *someValue)
{
    while(1==1) {
        sleep(1);
        printf("function 1 \n");
    }
}

void function2()
{
    while(1==1) {
        sleep(2);
        printf("function 2\n");
    }
}

int main()
{
    pthread_t thread_id;
    printf("I am inside the main function\n");
    pthread_create(&thread_id, NULL, function1, NULL);
    function2();
    exit(0);
}
```



# compile & run

```
(base) mahmutunan@MacBook-Pro lecture31 % gcc exercise2.c -o exercise2 -lpthread
(base) mahmutunan@MacBook-Pro lecture31 % ./exercise2
I am inside the main function
function 1
function 2
function 1
function 1
function 2
function 1
function 1
function 2
function 1
function 1
function 2
function 1
function 1
```

# exercise 3

```
int globalVar = 50; //define a global variable

void *someFuncToCreateThread(void *someValue)
{
    int *threadId = (int *)someValue; // Store the value argument passed to this thread

    //define a static and a local variable
    static int staticVar = 75;
    int localVar = 10;

    // let's change the variables
    globalVar +=100;
    staticVar +=100;
    localVar +=100;
    printf("id =%d,global = %d,  local = %d, static =%d, \n", *threadId, globalVar,localVar ,staticVar);

    return NULL;
}

int main()
{
    int i;
    pthread_t thread_id;
    for (i = 0; i < 4; i++)
        pthread_create(&thread_id, NULL, someFuncToCreateThread, (void *)&thread_id);
    pthread_exit(NULL);
}
```

# compile & run

```
(base) mahmutunan@MacBook-Pro lecture31 % gcc exercise3.c -o exercise3 -lpthread
(base) mahmutunan@MacBook-Pro lecture31 % ./exercise3
id =151261184,global = 150,  local = 110, static =175,
id =151261184,global = 150,  local = 110, static =175,
id =151261184,global = 250,  local = 110, static =275,
id =151261184,global = 250,  local = 110, static =275,
(base) mahmutunan@MacBook-Pro lecture31 %
```

Remember, global and static variables are stored in data segment.

All threads share data segment, so they are shared by all threads.

# pthread1.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <pthread.h>
4
5  int nthreads;
6
7  void *compute(void *arg) {
8      long tid = (long)arg;
9
10     printf("Hello, I am thread %ld of %d\n", tid, nthreads);
11
12     return (NULL);
13 }
14
15 int main(int argc, char **argv) {
16     long i;
17     pthread_t *tid;
18
19     if (argc != 2) {
20         printf("Usage: %s <# of threads>\n", argv[0]);
21         exit(-1);
22     }
23
24     nthreads = atoi(argv[1]); // no. of threads
```

```

25
26 // allocate vector and initialize
27 tid = (pthread_t *)malloc(sizeof(pthread_t)*nthreads);
28
29 // create threads
30 for ( i = 0; i < nthreads; i++)
31     pthread_create(&tid[i], NULL, compute, (void *)i);
32
33 // wait for them to complete
34 for ( i = 0; i < nthreads; i++)
35     pthread_join(tid[i], NULL);
36
37 printf("Exiting main program\n");
38
39 return 0;
40 }
41

```

```

(base) mahmutunan@MacBook-Pro lecture31 % gcc pthread1.c -o exercise4 -lpthread
(base) mahmutunan@MacBook-Pro lecture31 % ./exercise4 4
Hello, I am thread 0 of 4
Hello, I am thread 1 of 4
Hello, I am thread 2 of 4
Hello, I am thread 3 of 4
Exiting main program

```

# pthread2.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <pthread.h>
4
5  int nthreads;
6
7  void *compute(void *arg) {
8      long tid = (long)arg;
9      pthread_t pthread_id = pthread_self();
10
11      printf("Hello, I am thread %ld of %d, pthread_self() = %lu (0x%lx)\n",
12            tid, nthreads, (unsigned long)pthread_id, (unsigned long)pthread_id);
13
14      return (NULL);
15 }
16
17 int main(int argc, char **argv) {
18     long i;
19     pthread_t *tid;
20     pthread_t pthread_id = pthread_self();
```

```

21
22 if (argc != 2) {
23     printf("Usage: %s <# of threads>\n", argv[0]);
24     exit(-1);
25 }
26
27 nthreads = atoi(argv[1]); // no. of threads
28
29 // allocate vector and initialize
30 tid = (pthread_t *)malloc(sizeof(pthread_t)*nthreads);
31
32 // create threads
33 for ( i = 0; i < nthreads; i++)
34     pthread_create(&tid[i], NULL, compute, (void *)i);
35
36 for ( i = 0; i < nthreads; i++)
37     printf("tid[%ld] = %lu (0x%lx)\n", i, tid[i], tid[i]);
38
39 printf("Hello, I am main thread. pthread_self() = %lu (0x%lx)\n",
40        (unsigned long)pthread_id, (unsigned long)pthread_id);
41
42 // wait for them to complete
43 for ( i = 0; i < nthreads; i++)
44     pthread_join(tid[i], NULL);
45
46 printf("Exiting main program\n");
47
48 return 0;
49 }

```

```
(base) mahmutunan@MacBook-Pro lecture31 % ./exercise5 4
tid[0] = 123145541038080 (0x70000e3ab000)
tid[1] = 123145541574656 (0x70000e42e000)
tid[2] = 123145542111232 (0x70000e4b1000)
tid[3] = 123145542647808 (0x70000e534000)
Hello, I am main thread. pthread_self() = 4365594048 (0x10435adc0)
Hello, I am thread 1 of 4, pthread_self() = 123145541574656 (0x70000e42e000)
Hello, I am thread 2 of 4, pthread_self() = 123145542111232 (0x70000e4b1000)
Hello, I am thread 0 of 4, pthread_self() = 123145541038080 (0x70000e3ab000)
Hello, I am thread 3 of 4, pthread_self() = 123145542647808 (0x70000e534000)
Exiting main program
```



# pthread3.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <pthread.h>
4
5  typedef struct foo {
6      pthread_t ptid; /* thread id returned by pthread_create */
7      int tid; /* user managed thread id (0 through nthreads-1) */
8      int nthreads; /* total no. of threads created */
9  } F00;
10
11 void *compute(void *args) {
12     F00 *info = (F00 *)args;
13     printf("Hello, I am thread %d of %d\n", info->tid, info->nthreads);
14
15     return (NULL);
16 }
17
18 int main(int argc, char **argv) {
19     int i, nthreads;
20     F00 *info;
21
22     if (argc != 2) {
23         printf("Usage: %s <# of threads>\n", argv[0]);
24         exit(-1);
25     }
```

```

27     nthreads = atoi(argv[1]); // no. of threads
28
29     // allocate structure
30     info = (F00 *)malloc(sizeof(F00)*nthreads);
31
32     // create threads
33     for ( i = 0; i < nthreads; i++) {
34         info[i].tid = i;
35         info[i].nthreads = nthreads;
36         pthread_create(&info[i].ptid, NULL, compute, (void *)&info[i]);
37     }
38
39     // wait for them to complete
40     for ( i = 0; i < nthreads; i++)
41         pthread_join(info[i].ptid, NULL);
42
43     free(info);
44     printf("Exiting main program\n");
45
46     return 0;
47 }

```

(base) mahmutunan@MacBook-Pro lecture31 % gcc pthread3.c -o exercise6 -lpthread

(base) mahmutunan@MacBook-Pro lecture31 % ./exercise6 4

Hello, I am thread 1 of 4

Hello, I am thread 0 of 4

Hello, I am thread 2 of 4

Hello, I am thread 3 of 4

Exiting main program